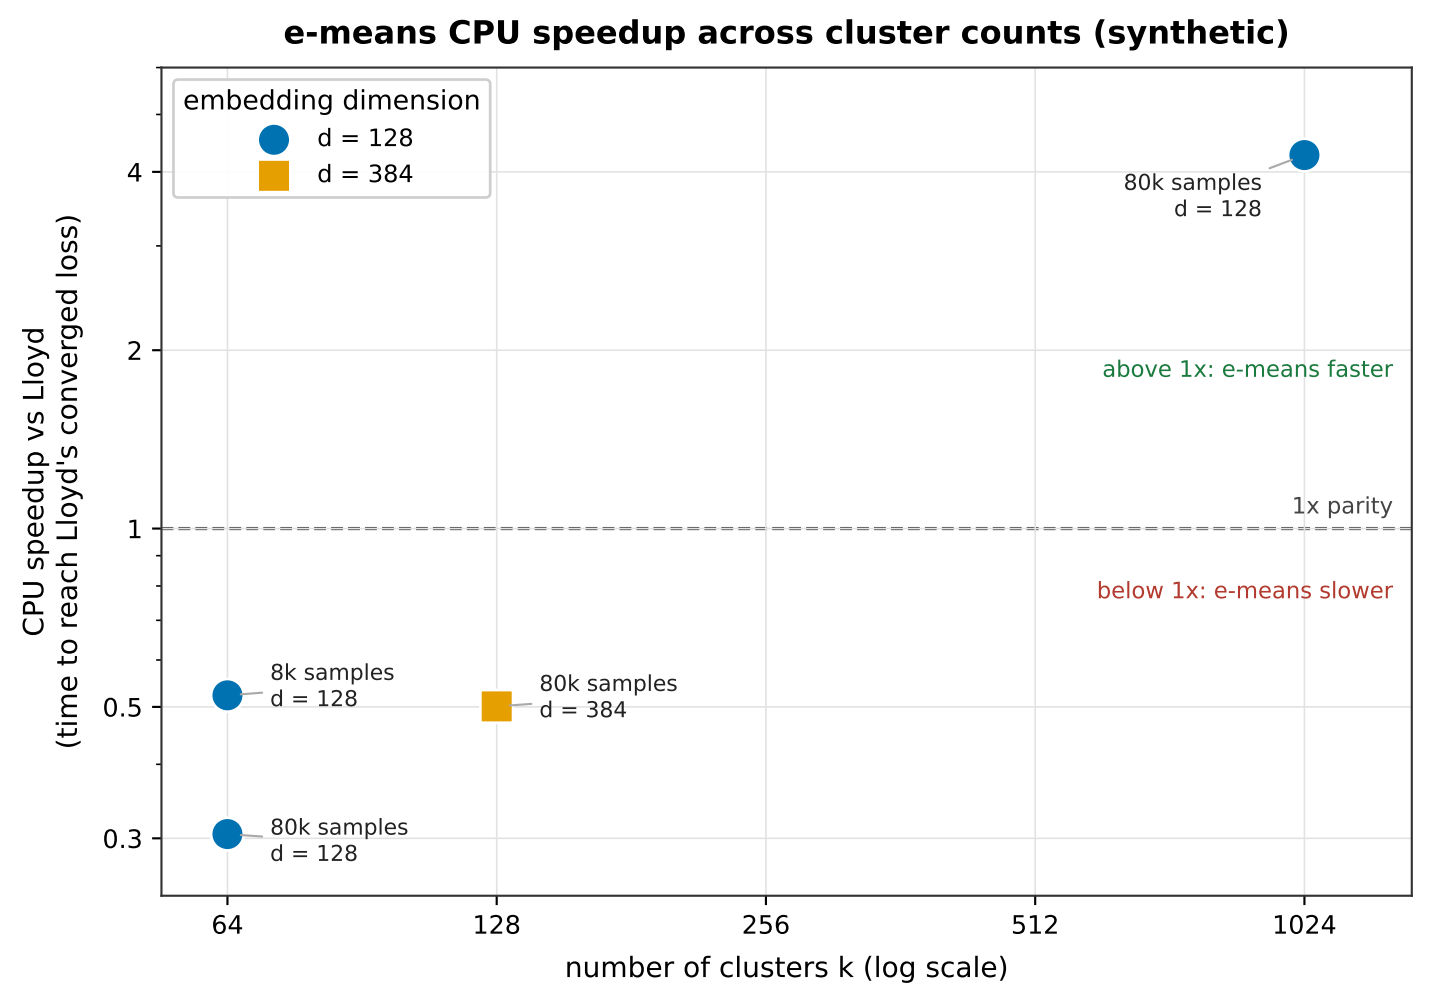
e-means CPU speedup across cluster counts (synthetic)
embedding dimension
d = 128
d = 384
80k samples
d = 128
above 1x: e-means faster
1x parity
below 1x: e-means slower
8k samples
d = 128
80k samples
d = 384
80k samples
d = 128
CPU speedup vs Lloyd
(time to reach Lloyd's converged loss)
4
2
1
0.5
0.3
64
128
256
512
1024
number of clusters k (log scale)